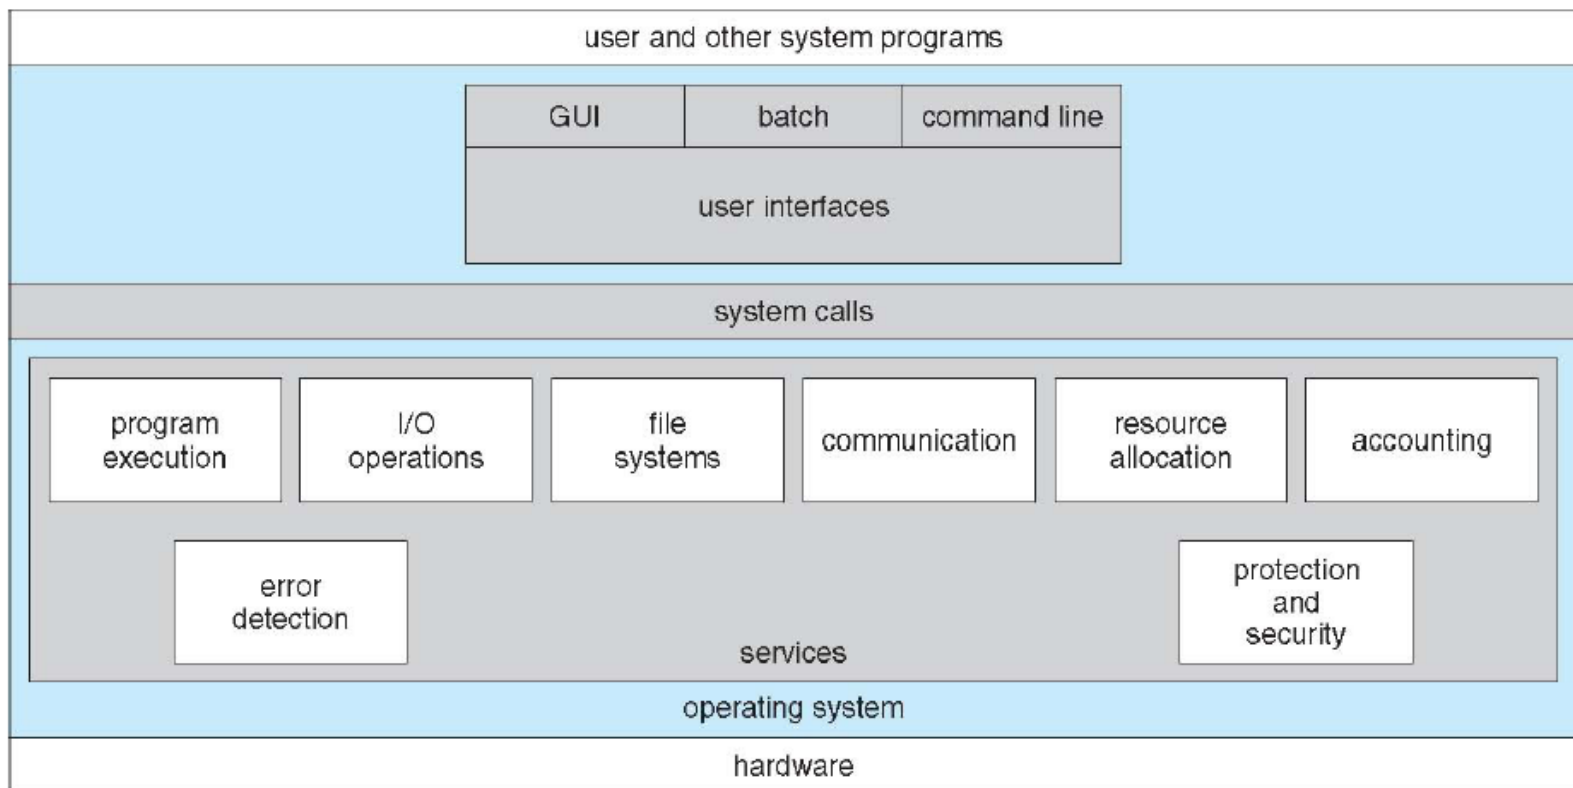


CH 2: Operating System Structure

Operating System Services

- A view of the Operating System services
 - An operating system provides an environment for the execution of programs.
 - It provides certain services to programs and to the users of those programs.



OS Services for User Support

- User Interface
 - CLI: Command-line interface
 - GUI: Graphical user interface
 - batch interface
- Program Execution
 - Load, run and end execution
- I/O operations
- File-system manipulation
 - Read/write/create/delete/search/list files
- Communications
 - shared memory vs. message passing
- Error Detection
 - Errors at the hardware-level or software-level
 - Bit-flip, memory error, power failure, parity, connection failure ... etc.

OS Services for Efficient Operation

■ Resource Allocation

- Resource: CPU cycles, memory space, memory bandwidth, file system space, file I/O bandwidth

■ Accounting

- Keeping track of which user (or process) uses how much of the resources

■ Protection and Security

- In multi user environment, all access to system resources must be controlled
- Other processes should not interfere with other processes

"A system is as strong as its weakest link"

OS Interfaces

- Command Interpreters (Command-line Interface: CLI)
 - Also called as shell: sh, ksh, csh, bash

Example) `$ rm file.txt`


The shell searches for 'rm' in the path, load them into memory and execute it with 'file.txt' as a parameter

- GUI (Graphical User Interface)
 - mouse, folder, icons, touch screen ... etc.

System Calls

■ Example of how system calls are used:

- Read data from one file and copy to another file
 - Request user for two file names
 - prompt use to enter
 - » each character being typed is displayed
 - In GUI, icons are shown and user selects the source file
 - Open two files
 - On error, error messages are printed and it terminates
 - Errors: file does not exist, permission error
 - ...
 - Copy data and write to the destination file in a loop
 - read and write
 - Close both files
 - Write completion message
 - Terminate the program



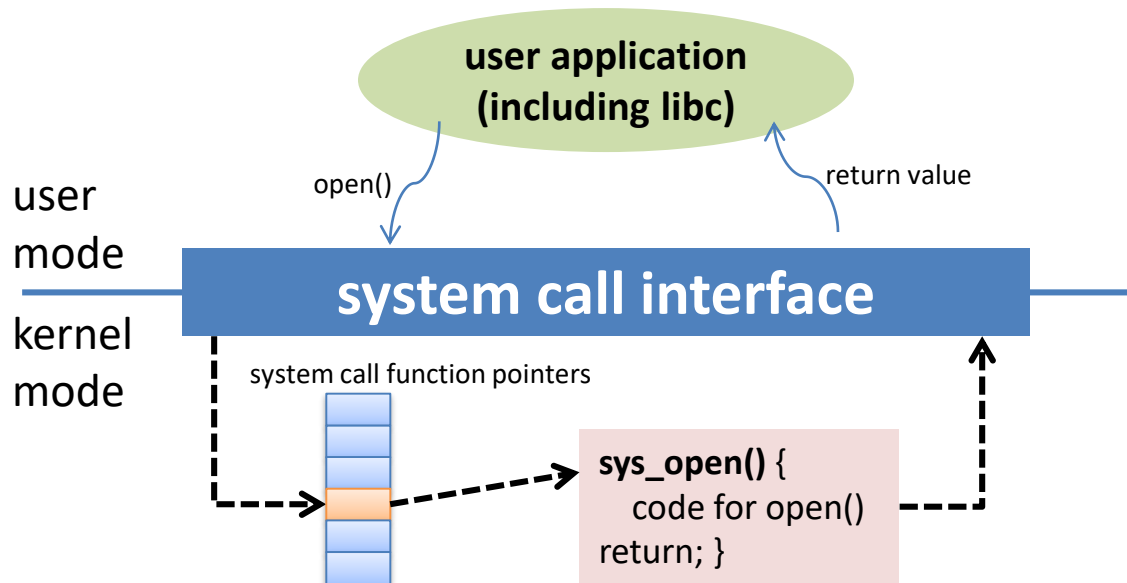
System Calls

- System calls are heavily used *even for a simple task*
 - thousands of system calls per second
- Programmers use system calls *indirectly* through APIs (Application Programming Interface: **set of functions**)
 - Windows API, POSIX API, Java API
 - API on top of another API
 - Java uses libc
- Why use APIs rather than system calls directly?
 - Directly using system call is difficult because ...
 - Portability (What is this?)

libc	system call
printf	write
write	
fread	read
read	
malloc	brk
pthread_lock	futex

System Call Interface

- Set of library functions that links to the system calls



- Caller does not need to know how system call is implemented
- Caller needs to know only the interface and what it returns

NAME

read - read from a file descriptor

SYNOPSIS

```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buf, size_t count);
```


System Call Interface

#	Name	Registers				Definition
		eax	ebx	ecx	edx	
0	sys_restart_syscall	0x00	-	-	-	kernel/signal.c:2058
1	sys_exit	0x01	int error_code	-	-	kernel/exit.c:1046
2	sys_fork	0x02	struct pt_regs *	-	-	arch/alpha/kernel/entry.S:716
3	sys_read	0x03	unsigned int fd	char __user *buf	size_t count	fs/read_write.c:391
4	sys_write	0x04	unsigned int fd	const char __user *buf	size_t count	fs/read_write.c:408
5	sys_open	0x05	const char __user *filename	int flags	int mode	fs/open.c:900
6	sys_close	0x06	unsigned int fd	-	-	fs/open.c:969
7	sys_waitpid	0x07	pid_t pid	int __user *stat_addr	int options	kernel/exit.c:1771
8	sys_creat	0x08	const char __user *pathname	int mode	-	fs/open.c:933
9	sys_link	0x09	const char __user *oldname	const char __user *newname	-	fs/namei.c:2520
10	sys_unlink	0x0a	const char __user *pathname	-	-	fs/namei.c:2352
11	sys_execve	0x0b	char __user *	char __user * __user *	char __user * __user *	arch/alpha/kernel/entry.S:925
12	sys_chdir	0x0c	const char __user *filename	-	-	fs/open.c:361
13	sys_time	0x0d	time_t __user *tloc	-	-	kernel/posix-timers.c:855
14	sys_mknod	0x0e	const char __user *filename	int mode	unsigned dev	fs/namei.c:2067
15	sys_chmod	0x0f	const char __user *filename	mode_t mode	-	fs/open.c:507
16	sys_lchown16	0x10	const char __user *filename	old_uid_t user	old_gid_t group	kernel/uid16.c:27
17	not implemented	0x11	-	-	-	
18	sys_stat	0x12	char __user *filename	struct __old_kernel_stat __user *statbuf	-	fs/stat.c:150
19	sys_lseek	0x13	unsigned int fd	off_t offset	unsigned int origin	fs/read_write.c:167
20	sys_getpid	0x14	-	-	-	kernel/timer.c:1337

System Call Types

- Six categories
 - Process control
 - **fork, exec, exit ...**
 - File manipulation
 - **create, open, close, read, write, lseek**
 - Device manipulation
 - **open, close, read, write, ioctl**
 - Information maintenance
 - **time, date, pid**
 - Communications
 - **open, close, connect, accept, read, write, send, recv, pipe, mmap, sendfile ...**
 - Protection
 - **Chmod, umask, chown ...**

System call types: Process control

- Type of tasks
 - end, abort
 - load, execute
 - create or terminate process
 - get/set process attributes
 - wait for time
 - wait event, signal event
 - allocate and free memory
- Scenarios
 - Error handling
 - Needs to terminate → creating a memory dump → print message → run next command
 - In GUI, pop-up window might be shown
 - In batch system, whole job is aborted
 - Loading and executing another program
 - Where to return the control to when new process terminates
 - Whether to concurrently run
 - In multiprogramming, we need to control processes
 - Get/set process attributes (e.g., priority, time limit)
 - Selectively terminate unneeded processes

System call types: Communications

- Two models of communications
 - *message passing* and *shared-memory* model
- Message passing
 - based on exchanging messages
 - connection must be first opened
 - What does the 'opening connection' really mean?
 - Sender/Receiver identities must be known
- Shared-memory model
 - In normal condition, processes are not allowed to access other processes' memory
 - System call for sharing must be invoked
 - Process must be careful to maintain the integrity of data in the shared memory region

Other types

- File management
 - create/delete file
 - open/close
 - read/write/reposition
 - get/set file attributes
 - file name, type, permissions, time, uid ... etc.
- Device management
 - request/release device: for exclusive usage \leftrightarrow open/close of files
 - read/write/reposition
 - get/set device attributes
 - logically attach/detach devices
- Information maintenance
 - get/set time or date
 - get/set system data: list of users, version info, free mem, disk usage
 - get/set process, file, device attributes
 - process execution time
 - dump, trace

OS Design and Implementation

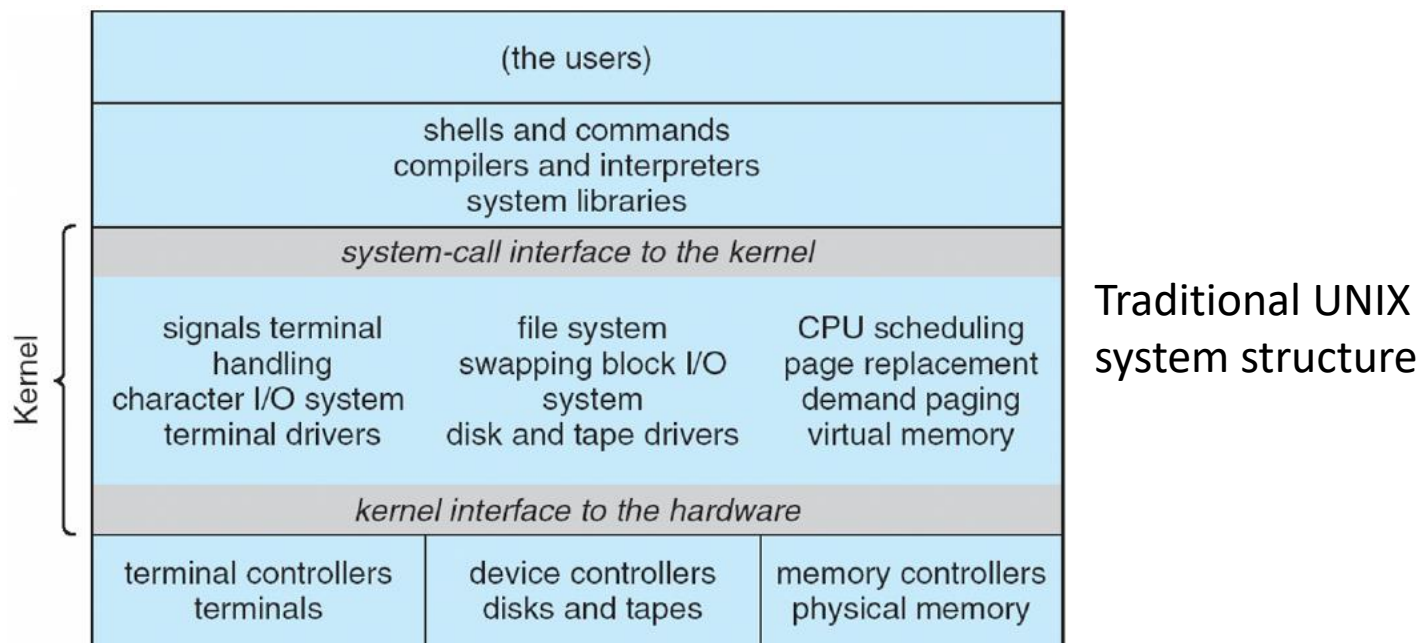
- Design goals: Requirements can be divided into two
 - User goals
 - ease of use, reliable, ease of learning, fast
 - System goals
 - ease of implementation, ease of maintenance, ease of operation, flexible, efficient, fault-tolerant
- **Vague requirements and no clear solution**
- Guideline: Mechanism and Policy
 - Separation of policy from mechanisms
 - mechanisms: how
 - policy: what
 - It is a technique frequently applied to many cases
 - ex) timer mechanisms and scheduling policy
 - ex) password expires in 30 days
 - Why separate them? Policy changes.
 - If policy change requires mechanisms change, it is not desirable.

OS Implementation

- Implementing in assembly language
 - tedious, difficult, error-prone, difficult to port
- Implementing in high-level language
 - Advantages
 - Faster implementation time
 - Compact and easier to understand/debug
 - Improved compiler will improve the generated code
 - Easier to port
 - MS-DOS (written in asm) needs emulator
 - Linux is mostly in C, no need to emulate, run natively
 - Disadvantages
 - Performance: potentially sub-optimal code generation
 - Not true in today's compiler technology
 - compiler is better at handling complex dependencies
 - Compromise
 - small critical code can be first written in high-level language and later replaced with the optimized ones

Operating System Structure

- **Monolithic** structure vs. componentized, modularized approach
 - component(module): well-defined portion of system with inputs, outputs and functions
- Simple Structure
 - Started simple and grew to become out-of-control
 - MS-DOS: applications can directly access I/O
 - Original UNIX (somewhat layered)



OS Structure: Layered Approach

■ Layered approach


- OS is broken into a number of layers
 - bottom layer: layer 0 = hardware
 - Layer M = user interface
- Layer M invokes operations of layer M-1
- Layer M provides operations that can be called by layer M+1

■ Advantage

- simplicity → easy to debug

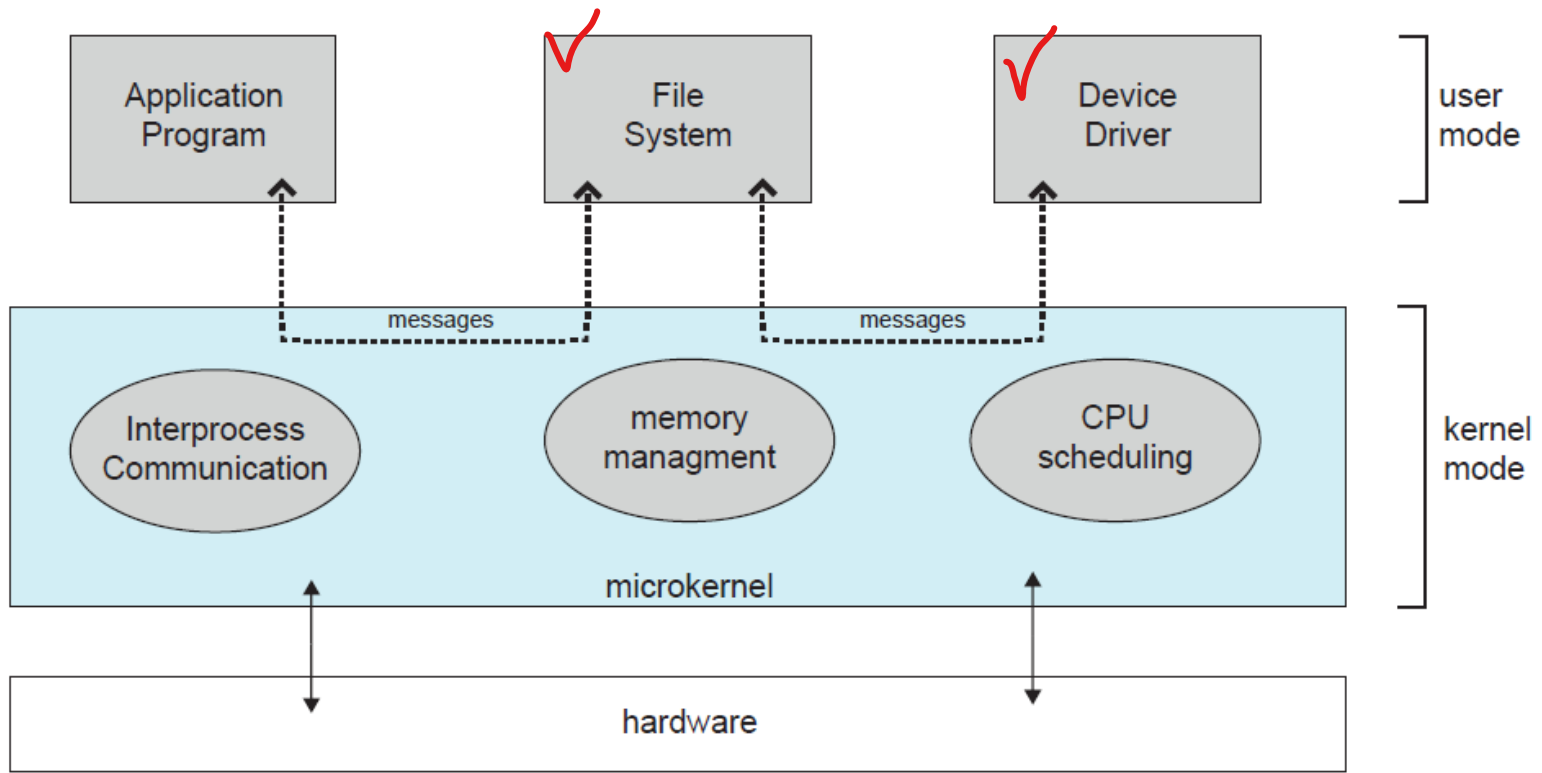
■ Disadvantage

- Defining layers is difficult
- Less efficient
 - crossing layers may add overhead



Microkernels

- Key idea
 - Make the kernel minimal
 - Implement non-essential functions as user-level library
- Main function
 - Providing efficient communication between applications and services
- Advantage
 - highly extensible OS kernel
 - easy to port kernels
 - more secure and reliable
- Disadvantage
 - Impact on performance





Modules

- Loadable kernel modules
 - A way to extend kernel functionality dynamically
 - A module that can be loaded dynamically into kernel
 - Used to add support for new H/W (device drivers), new file systems or new system calls
- Without LKM, kernel has to be rebuilt every time new feature is added
- Kernel has to include everything
 - Larger footprint
 - Larger memory consumption
- In Linux,
 - modules have *.ko extension
 - modprobe: add/remove modules
 - lsmod: show the status of modules
 - modinfo: show module information



Modules

■ lsmod

Module	Size	Used by
rfcomm	69632	2
xt_CHECKSUM	16384	1
iptable_mangle	16384	1
ipt_MASQUERADE	16384	4
nf_nat_masquerade_ipv4	16384	1 ipt_MASQUERADE
ipt_REJECT	16384	2
nf_reject_ipv4	16384	1 ipt_REJECT
xt_tcpudp	16384	6
xfrm_user	32768	1
xfrm_algo	16384	1 xfrm_user
iptable_nat	16384	1
nf_conntrack_ipv4	16384	3
nf_defrag_ipv4	16384	1 nf_conntrack_ipv4
nf_nat_ipv4	16384	1 iptable_nat
xt_addrtype	16384	2
ebtable_filter	16384	0
ebtables	36864	1 ebtable_filter
xt_conntrack	16384	2
nf_nat	24576	2 nf_nat_ipv4, nf_nat_masquerade_ipv4
br_netfilter	24576	0

■ modinfo

```
~$ modinfo nf_nat_ipv4
filename:      /lib/modules/4.4.0-66-generic/kernel/net/ipv4/netfilter/nf_nat_ipv4.ko
alias:         nf-nat-2
license:       GPL
srcversion:    62E44ADA1A1F6F29B5E6B7A
depends:        nf_nat, nf_conntrack
intree:        Y
vermagic:      4.4.0-66-generic SMP mod_unload modversions
```



OS Debugging

- Problem types
 - Fail-stop problem
 - Performance problem
- Failure analysis
 - Applications: log files
 - OS kernel: core (memory) dump
- Performance Tuning
 - Identify and remove the bottleneck